Designing an Index for ZooDB

Jonas Nick & Bogdan Vancea

May 29, 2014

Outline

ZooDB

- an open source object database written in Java
- JDO standard compliant
- 4 times faster than competitor db4o
- zoodb.org

Database Index

Key-Value datastructure that allows ordered iteration.

Database Index

Key-Value datastructure that allows ordered iteration.

Example:
```
ZooJdoHelper.createIndex(pm, Person.class, "name",
false);
```

Database Index

Key-Value datastructure that allows ordered iteration.

Example:
```
ZooJdoHelper.createIndex(pm, Person.class, "name",
false);
```

> Attribute
> Value $\rightarrow$ Object-ID

Database Index

Key-Value datastructure that allows ordered iteration.

Example:
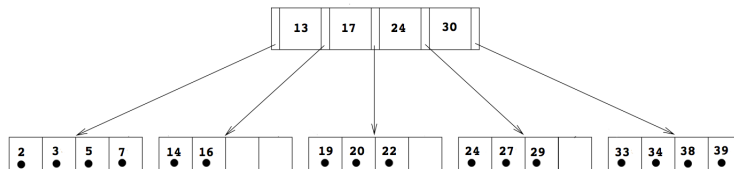ZooJdoHelper.createIndex(pm, Person.class, "name",
false);

```
Attribute
Value → Object-ID
```

```
OID
Object-ID        →
Diskpos
```
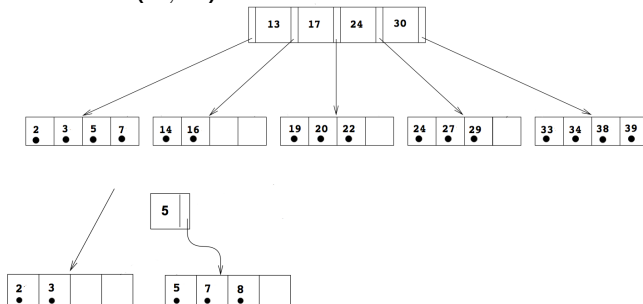
```
FSM
Page-ID → TxID
```

B+ Tree

- node fills one disk page
- inner node contains keys and children pointer, leaves contain keys and values
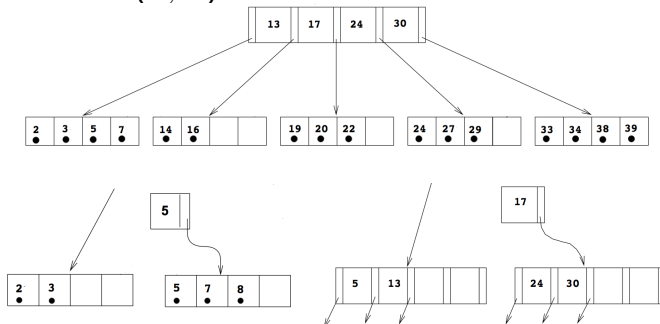- key unique vs. key-value unique

Example: insert $(8, v)$

Example: insert (8, *v*)

Example: insert (8, *v*)

Example: insert (8, *v*)

3908 loc + 2665 loc of test

Goals

- fast B+ tree index
- buffer manager to allow caching
- prefix sharing

Challenges - General

- edge cases
- runtime dominated by disk access
  - change nodes infrequently
  - fewer nodes is better
- only parent to child pointer
- determine best when its time to split/redistribute
- key unique vs. key-value unique
- buffer manager lookup takes time
- prefix-sharing encoding/decoding takes time
- prefix-sharing rebalancing takes time
- general optimizations
  - avoid polymorphism
  - bit-level operations

Prefix Sharing

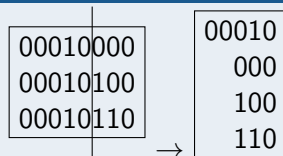## Exploit common prefix of keys in nodes

```
00010000
00010100
00010110
```

Prefix Sharing

Exploit common prefix of keys in nodes

```
00010000
00010100
00010110
              →
```
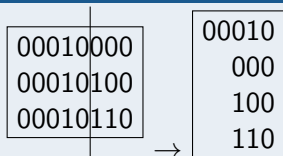
Prefix Sharing

Exploit common prefix of keys in nodes

Prefix Sharing

Exploit common prefix of keys in nodes

- variable number of key-value entries per page
- prefix determines
  - if 2 nodes can be merged without overflow
  - the number of entries that can be redistributed from one node to the other

Our B+ Tree

- 64 bit keys and 64 bit values
- key unique and key non-unique variants
- keys encoded based on common bit prefix

Class Diagram

Make it and add it here

Operations

- Search - Similar to normal B+ Tree
- Insert overflow
    - attempt to redistribute values to left sibling before creating a new node
- Delete underflow
    - check if possible to merge with left or right neighbour
    - check if possible to split current node between left and right
    - redistribute from left or right

Fine grained

- insert, remove, write
- duration, number of nodes
- prefix-sharing vs. no prefix-sharing

Whole system

- test harness
- partial PolePosition benchmark
- StackOverflow

Summary

- ...

Outlook

- ...