

Designing an Index for ZooDB

Jonas Nick & Bogdan Vancea

June 2, 2014

Outline

- 1 Introduction
- 2 Goals & Challenges
- 3 The new Index Implementation
- 4 Benchmarks



- ▶ an open source object database written in Java



- ▶ an open source object database written in Java
- ▶ based on JDO standard



- ▶ an open source object database written in Java
- ▶ based on JDO standard
- ▶ 4 times faster than competitor db4o



- ▶ an open source object database written in Java
- ▶ based on JDO standard
- ▶ 4 times faster than competitor db4o
- ▶ zoodb.org

Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

```
ZooJdoHelper.createIndex(pm, Person.class, "name",  
false);
```


Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

```
ZooJdoHelper.createIndex(pm, Person.class, "name",  
false);
```

Attribute Index
Value → Object-ID

Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

```
ZooJdoHelper.createIndex(pm, Person.class, "name",  
false);
```

Attribute Index
Value → Object-ID

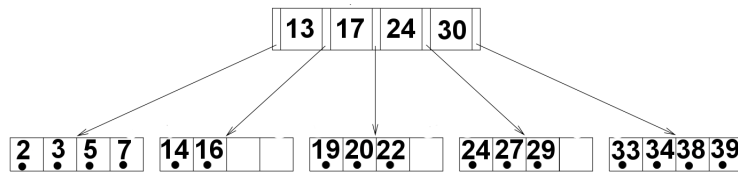
ObjectID Index
OID → Diskpos

Extension Index
Diskpos →
0|follow Diskpos

B+ Tree


Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

B+ Tree



- ▶ Inner node contains keys and children pointer, leaf contains keys and values.

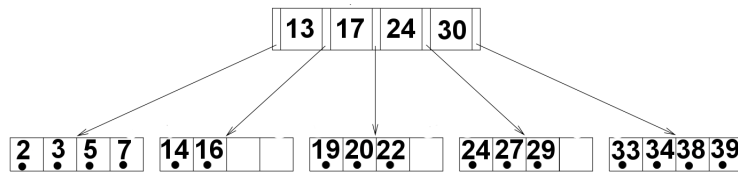
B+ Tree



- ▶ Inner node contains keys and children pointer, leaf contains keys and values.
- ▶ **Node fills one disk page.**

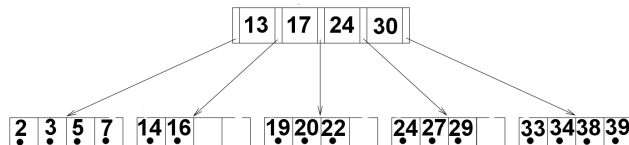
Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

B+ Tree

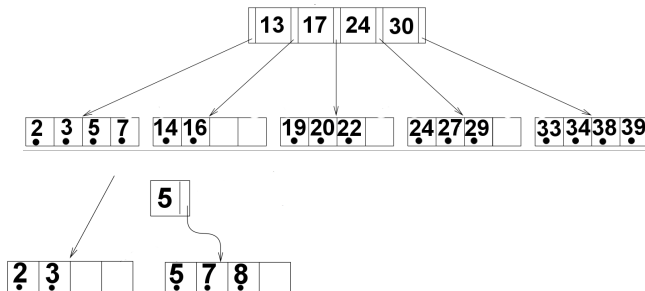


- ▶ Inner node contains keys and children pointer, leaf contains keys and values.
- ▶ **Node fills one disk page.**
- ▶ Node has maximum and minimum number of entries.

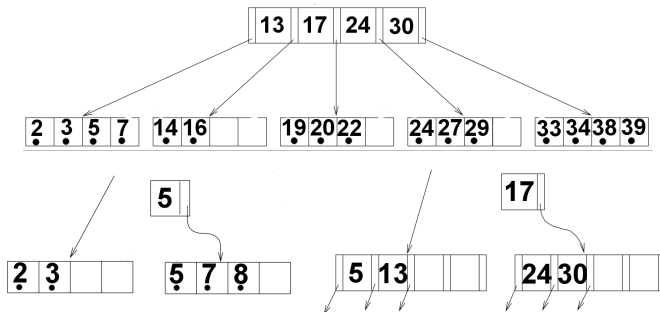
Example: insert (8, v)



Example: insert (8, v)

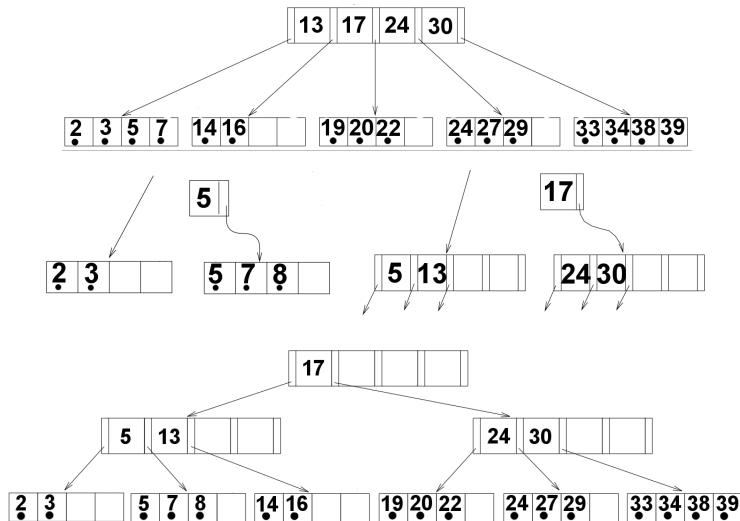


Example: insert (8, v)



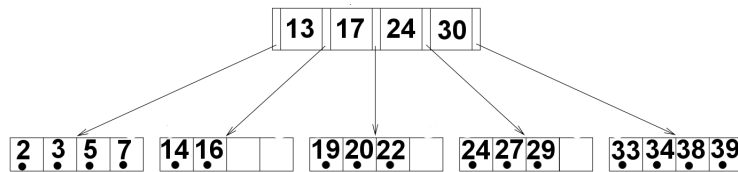
Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

Example: insert (8, v)



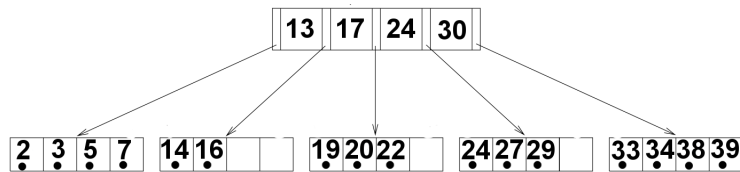
Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

B+ Tree




- ▶ Inner node contains keys and children pointer, leaf contain keys and values.
- ▶ **Node fills one disk page.**
- ▶ Node has maximum and minimum number of entries.

B+ Tree



- ▶ Inner node contains keys and children pointer, leaf contain keys and values.
- ▶ **Node fills one disk page.**
- ▶ Node has maximum and minimum number of entries.
- ▶ Rebalancing
 - ▶ on insert: split
 - ▶ on delete: merge or redistribute

B+ Tree



- ▶ Inner node contains keys and children pointer, leaf contain keys and values.
- ▶ **Node fills one disk page.**
- ▶ Node has maximum and minimum number of entries.
- ▶ Rebalancing
 - ▶ on insert: split
 - ▶ on delete: merge or redistribute
- ▶ Insert, remove, search are logarithmic.

Goals

- ▶ faster B+ tree index

Goals

- ▶ faster B+ tree index
- ▶ key unique and key-value unique
 - ▶ Ex. insert (1,1), (1,2)

Goals

- ▶ faster B+ tree index
- ▶ key unique and key-value unique
 - ▶ Ex. insert (1,1), (1,2)
- ▶ range query iterators

Goals

- ▶ faster B+ tree index
- ▶ key unique and key-value unique
 - ▶ Ex. insert (1,1), (1,2)
- ▶ range query iterators
- ▶ buffer manager to allow caching
 - ▶ fetches pages

Goals

- ▶ faster B+ tree index
- ▶ key unique and key-value unique
 - ▶ Ex. insert (1,1), (1,2)
- ▶ range query iterators
- ▶ buffer manager to allow caching
 - ▶ fetches pages
- ▶ prefix sharing

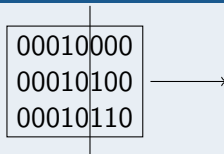
Prefix Sharing

Exploit common prefix

| |
|----------|
| 00010000 |
| 00010100 |
| 00010110 |


Prefix Sharing

Exploit common prefix



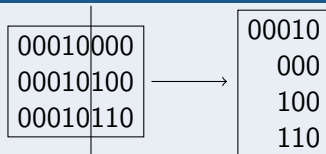
Prefix Sharing

Exploit common prefix



Prefix Sharing

Exploit common prefix



- ▶ allows storing more entries in a node
- ▶ determines if node under- or overflows

Challenges

- ▶ runtime dominated by disk access

Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes

Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes
 - ▶ rarely modify nodes

Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes
 - ▶ rarely modify nodes
- ▶ New features are costly.

Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes
 - ▶ rarely modify nodes
- ▶ New features are costly.
- ▶ Textbook algorithms need to be adapted.

Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes
 - ▶ rarely modify nodes
- ▶ New features are costly.
- ▶ Textbook algorithms need to be adapted.
 1. not optimized for our practical scenario

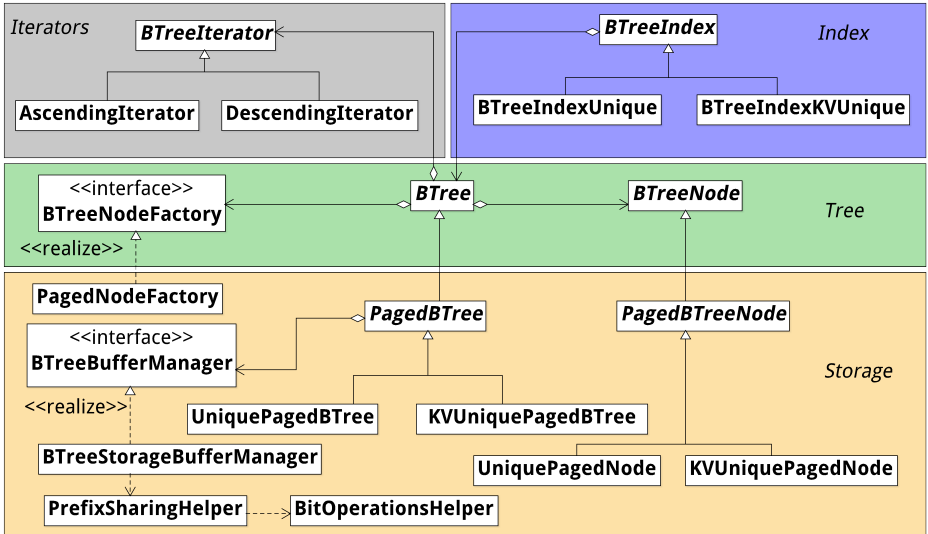
Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes
 - ▶ rarely modify nodes
- ▶ New features are costly.
- ▶ Textbook algorithms need to be adapted.
 1. not optimized for our practical scenario
 2. do not cover prefix sharing nor duplicates

Challenges

- ▶ runtime dominated by disk access
 - ▶ prefer fewer nodes
 - ▶ rarely modify nodes
- ▶ New features are costly.
- ▶ Textbook algorithms need to be adapted.
 1. not optimized for our practical scenario
 2. do not cover prefix sharing nor duplicates
- ▶ low-level implementation optimizations

Index Implementation



Operations

Operations

- ▶ Search - similar to normal B+ Tree

Operations

- ▶ Search - similar to normal B+ Tree
- ▶ Insert overflow
 1. redistribute left ?
 2. split

Operations

- ▶ Search - similar to normal B+ Tree
- ▶ Insert overflow
 1. redistribute left ?
 2. split
- ▶ Delete underflow
 1. merge with left/right ?
 2. split between left and right ?
 3. redistribute left/right

Operations

- ▶ Search - similar to normal B+ Tree
- ▶ Insert overflow
 1. redistribute left ?
 2. split
- ▶ Delete underflow
 1. merge with left/right ?
 2. split between left and right ?
 3. redistribute left/right
- ▶ Write
 - ▶ only write dirty nodes
 - ▶ prefix encoding

Operations

- ▶ Search - similar to normal B+ Tree
- ▶ Insert overflow
 1. redistribute left ?
 2. split
- ▶ Delete underflow
 1. merge with left/right ?
 2. split between left and right ?
 3. redistribute left/right
- ▶ Write
 - ▶ only write dirty nodes
 - ▶ prefix encoding
- ▶ insert/delete more costly, exactly how much?

Microbenchmarks

- ▶ full in-memory, index only tests

Microbenchmarks

- ▶ full in-memory, index only tests

Duration - Old index is the baseline

| Operation | No Prefix sharing | Prefix sharing |
|-----------|-------------------|----------------|
| Search | 1 | 0.9 - 1.1 |
| Insert | 1 | 1.6 - 2.8 |
| Delete | 1 | 1.45 - 2.9 |

Microbenchmarks

- ▶ full in-memory, index only tests

Duration - Old index is the baseline

| Operation | No Prefix sharing | Prefix sharing |
|-----------|-------------------|----------------|
| Search | 1 | 0.9 - 1.1 |
| Insert | 1 | 1.6 - 2.8 |
| Delete | 1 | 1.45 - 2.9 |

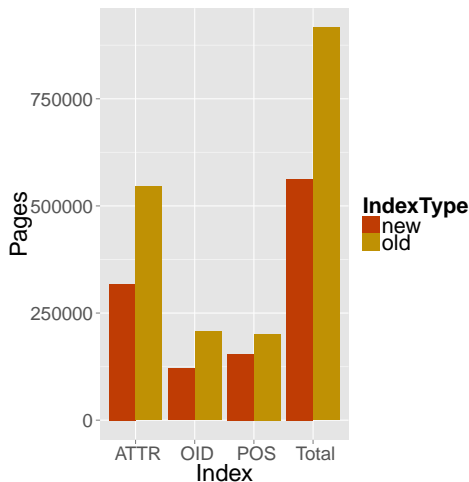
Size of B+ tree - Old index is the baseline

| Operation | No Prefix sharing | Prefix sharing |
|-----------|-------------------|----------------|
| Insert | 1 | 0.5 - 1.1 |
| Delete | 1 | 0.5 - 0.75 |

StackOverflow Data Import

- ▶ real-world workload
- ▶ StackOverflow data
 - ▶ 1.3 million users
 - ▶ 10.3 million posts
 - ▶ 13 million comments
 - ▶ 25 million votes
- ▶ 3 key unique attribute indexes
- ▶ 9 key-value unique attribute indexes


StackOverflow Import - Index Sizes



- ▶ page size: 4KB
- ▶ database size: 31 GB


| Index | Space saving (%) |
|-----------|------------------|
| Attribute | 41.6 |
| OID | 41.5 |
| POS | 23.1 |
| Total | 38.5 |

StackOverflow Import - Commit times



- ▶ import with new index 25% faster
- ▶ why?

StackOverflow Import - Commit times



- ▶ import with new index 25% faster
- ▶ why?
- ▶ more entries in a node
→ fewer dirty nodes
- ▶ data locality

Summary

- ▶ prefix sharing: trade-off between speed and space
- ▶ works well in practice
- ▶ microbenchmarks
- ▶ implementation complexity.

Q&A

- ▶ Thank you for your attention!
- ▶ Questions ?