

Designing an Index for ZooDB

Jonas Nick & Bogdan Vancea

June 2, 2014

Outline

- 1 Introduction
- 2 Goals & Challenges
- 3 The new Index Implementation
- 4 Benchmarks



- an open source object database written in Java



- an open source object database written in Java
- based on JDO standard



- an open source object database written in Java
- based on JDO standard
- 4 times faster than competitor db4o



- an open source object database written in Java
- based on JDO standard
- 4 times faster than competitor db4o
- zoodb.org

Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

```
ZooJdoHelper.createIndex(pm, Person.class, "name",  
false);
```


Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

```
ZooJdoHelper.createIndex(pm, Person.class, "name",  
false);
```

Attribute Index

Value → Object-ID

Database Index

Key-Value data structure

1. **fast** retrieval
2. **ordered** iteration
3. stored in a **file**

```
ZooJdoHelper.createIndex(pm, Person.class, "name",  
false);
```

Attribute Index
Value → Object-ID


ObjectID Index
OID → Diskpos

Extension Index
Diskpos →
0|follow Diskpos

B+ Tree


Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

B+ Tree




- Inner node contains keys and children pointer, leaf contains keys and values.

B+ Tree




- Inner node contains keys and children pointer, leaf contains keys and values.
- **Node fills one disk page.**

B+ Tree




- Inner node contains keys and children pointer, leaf contains keys and values.
- **Node fills one disk page.**
- Node has maximum and minimum number of entries.

Example: insert (8, v)




Example: insert (8, v)



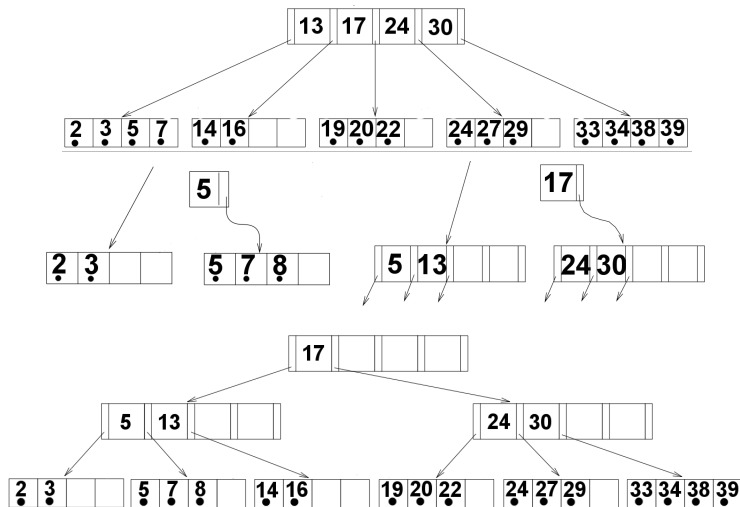
Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

Example: insert (8, v)




Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

Example: insert (8, v)




Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

B+ Tree



- Inner node contains keys and children pointer, leaf contain keys and values.
- **Node fills one disk page.**
- Node has maximum and minimum number of entries.


B+ Tree



- Inner node contains keys and children pointer, leaf contain keys and values.
- **Node fills one disk page.**
- Node has maximum and minimum number of entries.
- Rebalancing
 - on insert: split
 - on delete: merge or redistribute

Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

B+ Tree



- Inner node contains keys and children pointer, leaf contain keys and values.
- **Node fills one disk page.**
- Node has maximum and minimum number of entries.
- Rebalancing
 - on insert: split
 - on delete: merge or redistribute
- Insert, remove, search are logarithmic.

Images adapted from Database Management Systems by Ramakrishnan and Gehrke.

Goals

- faster B+ tree index

Goals

- faster B+ tree index
- key unique and key-value unique
 - Ex. insert (1,1), (1,2)

Goals

- faster B+ tree index
- key unique and key-value unique
 - Ex. insert (1,1), (1,2)
- range query iterators

Goals

- faster B+ tree index
- key unique and key-value unique
 - Ex. insert (1,1), (1,2)
- range query iterators
- buffer manager to allow caching
 - fetches pages

Goals

- faster B+ tree index
- key unique and key-value unique
 - Ex. insert (1,1), (1,2)
- range query iterators
- buffer manager to allow caching
 - fetches pages
- prefix sharing


Prefix Sharing

Exploit common prefix

00010000
00010100
00010110


Prefix Sharing

Exploit common prefix




Prefix Sharing

Exploit common prefix



Prefix Sharing

Exploit common prefix



- allows storing more entries in a node
- determines if node under- or overflows

Challenges

- runtime dominated by disk access

Challenges

- runtime dominated by disk access
 - prefer fewer nodes

Challenges

- runtime dominated by disk access
 - prefer fewer nodes
 - rarely modify nodes

Challenges

- runtime dominated by disk access
 - prefer fewer nodes
 - rarely modify nodes
- New features are costly.

Challenges

- runtime dominated by disk access
 - prefer fewer nodes
 - rarely modify nodes
- New features are costly.
- Textbook algorithms need to be adapted.

Challenges

- runtime dominated by disk access
 - prefer fewer nodes
 - rarely modify nodes
- New features are costly.
- Textbook algorithms need to be adapted.
 1. not optimized for our practical scenario

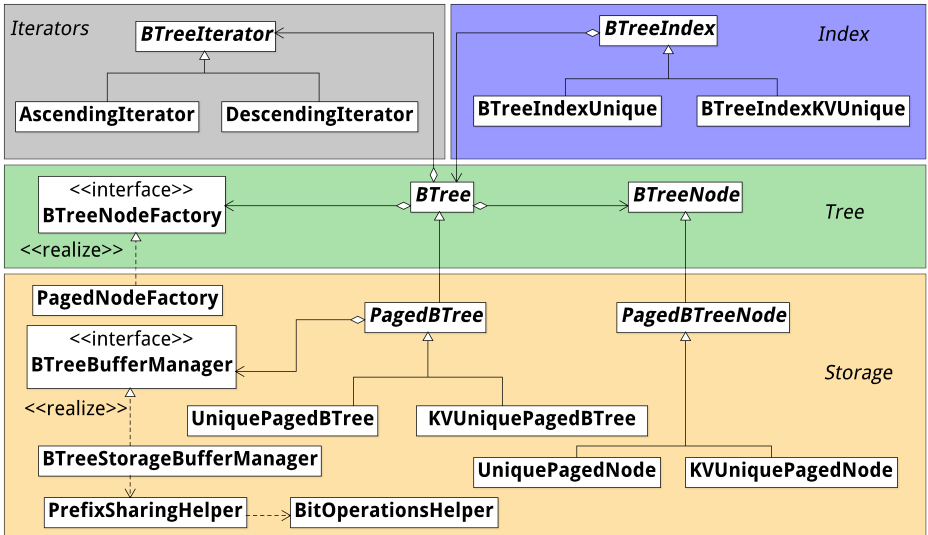
Challenges

- runtime dominated by disk access
 - prefer fewer nodes
 - rarely modify nodes
- New features are costly.
- Textbook algorithms need to be adapted.
 1. not optimized for our practical scenario
 2. do not cover prefix sharing nor duplicates

Challenges

- runtime dominated by disk access
 - prefer fewer nodes
 - rarely modify nodes
- New features are costly.
- Textbook algorithms need to be adapted.
 1. not optimized for our practical scenario
 2. do not cover prefix sharing nor duplicates
- low-level implementation optimizations

Index Implementation



Operations

Operations

- Search - similar to normal B+ Tree

Operations

- Search - similar to normal B+ Tree
- Insert overflow
 1. redistribute left ?
 2. split

Operations

- Search - similar to normal B+ Tree
- Insert overflow
 1. redistribute left ?
 2. split
- Delete underflow
 1. merge with left/right ?
 2. split between left and right ?
 3. redistribute left/right

Operations

- Search - similar to normal B+ Tree
- Insert overflow
 1. redistribute left ?
 2. split
- Delete underflow
 1. merge with left/right ?
 2. split between left and right ?
 3. redistribute left/right
- Write
 - only write dirty nodes
 - prefix encoding

Operations

- Search - similar to normal B+ Tree
- Insert overflow
 1. redistribute left ?
 2. split
- Delete underflow
 1. merge with left/right ?
 2. split between left and right ?
 3. redistribute left/right
- Write
 - only write dirty nodes
 - prefix encoding
- insert/delete more costly, exactly how much?

Microbenchmarks

- full in-memory, index only tests

Microbenchmarks

- full in-memory, index only tests

Duration - Old index is the baseline

Operation	No Prefix sharing	Prefix sharing
Search	1	0.9 - 1.1
Insert	1	1.6 - 2.8
Delete	1	1.45 - 2.9

Microbenchmarks

- full in-memory, index only tests

Duration - Old index is the baseline

Operation	No Prefix sharing	Prefix sharing
Search	1	0.9 - 1.1
Insert	1	1.6 - 2.8
Delete	1	1.45 - 2.9


Size of B+ tree - Old index is the baseline

Operation	No Prefix sharing	Prefix sharing
Insert	1	0.5 - 1.1
Delete	1	0.5 - 0.75

StackOverflow Data Import

- real-world workload
- StackOverflow data
 - 1.3 million users
 - 10.3 million posts
 - 13 million comments
 - 25 million votes
- 3 key unique attribute indexes
- 9 key-value unique attribute indexes


StackOverflow Import - Index Sizes



- page size: 4KB
- database size: 31 GB


Index	Space saving (%)
Attribute	41.6
OID	41.5
POS	23.1
Total	38.5

StackOverflow Import - Commit times



- import with new index 25% faster
- why?

StackOverflow Import - Commit times



- import with new index 25% faster
- why?
- more entries in a node
→ fewer dirty nodes
- data locality

Summary

- prefix sharing: trade-off between speed and space
- works well in practice
- microbenchmarks
- implementation complexity.

Q&A

- Thank you for your attention!
- Questions ?